

Critical Capabilities for Decision Intelligence Platforms

27 January 2026 - ID G00827620 - 52 min read

By Carlie Idoine, David Pidsley, [and 3 more](#)

Organizations must improve decision quality and outcomes with AI-first automation that is safe. Leaders responsible for decision intelligence should select vendors with critical capabilities for decision analysis, low-code engineering, decision science and decision governance.

Overview

Key Findings

- While some vendors provide decision intelligence platforms (DIPs) as their primary offering, many ecosystem players have expanded into the DIP market by incorporating DIP capabilities within their broader ecosystems.
- Users of DIPs and their support tools differ in their needs and requirements. Some are no-code decision analysts, others are low-code decision engineers and others still are code-first decision scientists.
- Two patterns for decision modeling have emerged in this market. Procedural decision modeling is the stepwise process of modeling decision flows, and declarative decision modeling is the decomposition and hierarchical modeling of decision policies.
- In the DIP market, AI agents, which are important in many workflows today, are part of the end-to-end insight-to-action workflow, or decision life cycle. For example, AI agents for DI, used in designing decision models, are viewed by DIP vendors as an alternative to providing decision blueprints and templates.

Recommendations

- Choose the DIP that best helps to model decisions and establish high-quality decision making that is coherent, contextual, connected, compliant, cost-optimized, collaborative and continuous across your organization.
- Emphasize persona-based use cases to prioritize skills and people when selecting the most appropriate DIP for your organization. Choose the no-code, low-code and code-first features that best fit your needs.
- Investigate existing solutions for AI techniques within your current technology landscape, and weigh them against the benefits of DIP vendors that provide a composable architecture and composite AI capabilities consistent with those needed for decision augmentation.
- Gauge the level of human-AI delegation appropriate for relevant decision services today and in the medium term, rather than rushing to decision automation or prematurely risking the use of autonomous AI agents in decision execution.

Strategic Planning Assumptions

- By 2027, 25% of ungoverned decisions using large language models (LLMs) will cause financial or reputational loss due to human biases, insufficient critical thinking, and AI sycophancy.
- By 2027, 50% of business decisions will have been augmented or automated by AI agents for decision intelligence.
- By 2028, 25% of CDAO vision statements will become “decision-centric,” surpassing “data-driven” slogans, with human decision-making behaviors explicitly addressed to improve D&A value.
- By 2030, explicitly modeled business decisions will be five times more trusted and 80% faster than ungoverned decisions, enabled by decision intelligence platform adoption.

What You Need to Know

Global competition, geopolitical disruptions and growing demands from customers are putting new pressures on organizations to clarify and accelerate their business decision-making processes. These scenarios are further complicated by the accelerated pace of digital business and new regulations regarding fairness and transparency. The pressure on data and analytics (D&A) leaders has never been greater.

While the need to establish high-quality decision making that is coherent, contextual, connected, compliant, cost-optimized, collaborative and continuous has become a priority for organizations, the ability to coordinate activities and respond efficiently and effectively is often hindered by disconnected processes. It is not enough to offer fragmented, data-driven insights through business dashboards or other siloed, general-purpose, non-contextualized analytical applications. The growing power of AI is offering organizations game-changing opportunities to bring decision making to a higher level of strategic and competitive differentiation and achieve a decision advantage. Many organizations, however, still need assistance to turn the promise of AI into reality and close the gap between envisioning the technological possibilities and actual business decision making.

Decision intelligence platforms (DIPs) can help D&A leaders address and fill critical needs by providing software for decision service development that supports, augments and automates decision making of humans or machines. All of this is powered by the composition of data, analytics, knowledge and AI techniques. The use cases for DIPs — which we define as decision analysis, decision engineering, decision science and decision stewardship — are reflective of the user personas and related decision intelligence skills needed to define, manage, observe and collaborate across the decision life cycle.


Analysis

Critical Capabilities Use-Case Graphics

Vendor Product Scores for the Decision Analysis Use Case



Product or Service Scores for Decision Analysis



As of 8 December 2025


1 2 3 4 5 © Gartner, Inc

Gartner

Vendor Product Scores for the Decision Engineering Use Case



Product or Service Scores for Decision Engineering




Gartner

Vendor Product Scores for the Decision Science Use Case



Product or Service Scores for Decision Science




Gartner



Vendor Product Scores for the Decision Stewardship Use Case

Product or Service Scores for Decision Stewardship



Gartner

Vendors

ACTICO

ACTICO offers its Decision Platform, which specializes in enabling regulated industries — such as financial services — to design, execute and govern decisions in environments where precision, compliance and trust are paramount. The platform comprises several integrated modules. These include ACTICO Modeler, which provides an intuitive graphical interface for rule creation, organization and test case management, and ACTICO Model Hub for decision model life cycle management. The modules also include

the ACTICO Platform CLI for loading artificial intelligence/machine learning (AI/ML) models into the ACTICO Model Hub. ACTICO Execution Server is available for hosting decision services as scalable web services and ACTICO Execution CLI for batch processing capabilities. ACTICO Engine can be embedded into Java applications via API. ACTICO Workplace enables rules-driven interactive decision applications that can be integrated into workflows and dashboards.

ACTICO Decision Platform offers strong capabilities for decision execution, decision governance and decision service composition, ensuring decisions are consistent and scalable with appropriate oversight and control. Decision execution is supported by extensive deployment versatility and comprehensive DecisionOps. Decision governance is enhanced by strong, explainable AI techniques targeted at both business and technical personas. Decision services composition is supported by decision services that provide version control and traceability. The low-code, composable architecture — in which decision logic and components are modular, autonomous, discoverable and reusable — provides not only flexibility but also enhanced usability. Additionally, the availability of dependency graphs that show relationships between services provides visibility into the complex relationships between the various decision components for more transparent, improved decision making.

Optimization and simulation techniques, graph and knowledge techniques, and ML techniques are areas for growth and enhancement. Optimization capabilities are not native to the platform, while simulation and scenario techniques are limited. Graph and knowledge techniques are provided by limited out-of-the-box capabilities and enhanced via third-party integrations.

ACTICO Decision Platform is best-suited for the decision engineering use case, thanks to its strong decision service composition capabilities. It is also suited for the decision stewardship use case due to its solid decision governance capabilities.

Aera Technology

Aera Decision Cloud is a strong platform for operationalizing decision intelligence at scale, with its unified approach to multimodal data connectivity, advanced ML and decision service composition within an API-first architecture. Enterprises benefit from integrated optimization, simulation and human-in-the-loop capabilities, especially in the supply chain domain.

Aera Decision Cloud is strong in multimodal D&A preparation, with the ability to extract and ingest data in-memory. It is also strong in ML and D&A preparation, and offers extensive out-of-the-box connectors for diverse data sources across on-premises and cloud environments. Seamless data exploration and ingestion is supported by a multiexperience UI that optimizes analytics development and content consumption for different decision-making processes. Users interact through designing workspaces, recommendations and an LLM interface, enhancing accessibility and usability. Aera Decision Cloud demonstrates strong automated ML capabilities, particularly in supervised learning, allowing users to build accurate predictive models from historical data. It supports both low-code and code-first development via notebooks and an advanced code editor. The platform's composite semantic layer and action framework expose decision data models as web services and facilitate human-in-the-loop decision recommendations that can trigger external systems and business process execution.

Aera Decision Cloud has opportunities for improvement in decision collaboration and AI agents for decision intelligence. Analytics and recommendations are not easily embedded within third-party applications, requiring decision makers and decision stewards to operate within the Aera environment. While its support for small language model (SLM) integration and generative AI (GenAI) shows promise, these capabilities are still developing. Deployment versatility is limited; although the platform is Kubernetes-based, there is no on-premises option. Additionally, the platform currently does not offer AI agents to support decision modeling, enhance collaboration or assist with decision service composition.

Aera provides modular software components that are easily consumed or composed by other applications, making Aera Decision Cloud suitable for decision engineers. Its well-documented API with SDKs and accessible platform objects support internal and external integration, catering to decision scientists with strong optimization and simulation features. With its strengths in ML, optimization and business intelligence (BI), Aera Decision Cloud is also well-suited for enterprises seeking agentic decision intelligence capabilities.

CRIF

The CRIF StrategyOne platform provides a low-code decision intelligence environment for configuring rules and decision flows, with integrated governance features that include auditability and compliance controls. These capabilities combine to provide easy

accessibility for nontechnical users and accelerated development and deployment of decision flows in a controlled and managed environment. The platform supports the design, optimization and execution of decision strategies across credit life cycle processes such as origination, risk-based pricing, portfolio management, fraud detection and collections. StrategyOne is available for deployment on-premises, in private or public cloud, or as a SaaS solution.

CRIF's strongest capabilities are in decision modeling, decision service composition, and rule and logic-based techniques. Its decision-centric UI enables intuitive, no-code modeling, allowing business users to easily and rapidly define decision logic. For decision service composition, CRIF provides robust capabilities to define, package and reuse decision models as modular decision services that can be defined once and then reused and leveraged across the platform. For rule- and logic-based techniques, it supports declarative rule expression — i.e., enabling business users to define and articulate rule logic — with accessible formats like decision tables and trees.

CRIF's biggest areas for improvement lie in its lack of native features for graph and knowledge techniques, AI agents for DI, multimodal D&A preparation, real-time event stream processing, and ML. In relation to graph and knowledge techniques, it lacks features for graph analytics, graph data science and knowledge graphs. AI agents for practicing DI techniques in any area of the decision life cycle — including modeling, execution or monitoring — were not in general availability. CRIF lacks native multimodal D&A preparation techniques, especially in data preparation. It lacks real-time event stream capabilities and does not support low-latency ingestion or complex event detection. The platform requires users to rely on other vendors' tools to access, combine and transform data, develop ML models, design composite AI-based decision flows or implement streaming decisions.

CRIF also has opportunities to enhance capabilities across all use cases, as it is weak in those capabilities most critical for decision analysts, decision engineers, decision scientists and decision stewards.

Decisions

Decisions is a low-code platform that combines workflow and business rule engines with visual designers for creating flows, forms, dashboards and reports. The platform supports modeling, testing and deployment of decision services, with execution options for real-time streaming and batch processes. Decisions offers a no-code, web-based design

studio for building and managing complex business logic. It includes governance features, such as versioning and auditing. Deployment options include SaaS, public/private cloud, hybrid or on-premises.

The strongest of Decisions' capabilities are real-time event stream processing, decision modeling, decision execution and decision collaboration, which help organizations productize high-value use cases as applications. For data streaming, Decisions offers significant support for low-latency ingestion (e.g., through messaging services like Kafka and RabbitMQ), complex event detection (e.g., with event stream listeners and temporal controls) and continuous intelligence (e.g., by allowing the rules and ML for analytics and notifications with streaming data). For decision modeling, the platform is strong in decision logic and mapping, decision blueprints and templates, and combined application (or fusion) of different AI techniques for composite AI. For decision execution, Decisions offers runtime decision orchestration, with features supporting scaling for real-time decisions and DecisionOps capabilities for managing the decision life cycle. For decision collaboration, the platform enables agile collaboration of decision flows with features for sprints and user stories, comment history, differential analysis, rollback functionality and built-in testing capabilities.

Customers evaluating Decisions' platform should be aware of its limited options for optimization and simulation, graph techniques and ML. The platform either did not demonstrate capabilities or offered limited features to support these techniques. In most cases, users would need to develop code or models on other platforms that could then be run on the Decisions platform, or connect to other platforms to access these capabilities.

Decisions' strongest use cases are decision engineering and decision stewardship. Its strength in decision engineering comes from its decision service composition, decision modeling and execution. For decision stewardship, the platform benefits from real-time event stream techniques, decision monitoring and decision collaboration.

Faculty

Faculty's Frontier decision intelligence platform is built around Frontier's Computational Twin framework. This framework connects data and AI assets to power coherent optimizations and simulations of the business and helps decision makers use those capabilities to close decision loops.

Faculty Frontier is strong in simulation techniques, decision monitoring and BI, helping organizations seeking to maximize benefits and manage trade-offs by identifying optimal operational and strategic decisions under constraints. Faculty Frontier's optimization and simulation tools empower users to conduct what-if analyses and scenario planning, simulating environments and outcomes to support decision making. These features are particularly valuable for complex operational challenges, strategic planning and resource allocation. Decision monitoring helps visualize decision points in context and how decisions affect metrics, and enables learning from human feedback loops to capture judgment or expertise. The platform is strong in BI and its composite semantic layer provides a consistent view of data, decisions and processes.

Potential growth and enhancement areas for Faculty are graph and knowledge techniques, real-time event stream processing and multimodal D&A preparation. It does not support graph data science, nor does it provide an underlying graph database that could inhibit discovery of relationships between entities and reduce the effectiveness of predictive models based on relationships. For real-time event stream techniques, Faculty Frontier's lack of complex event detection and continuous intelligence limits its ability to execute decisions that are informed by real-time context; however, it can support near-real-time, event-informed decisioning based on frequently updated data. A weakness in multimodal D&A preparation will make it hard for users to create analytic models, user-defined measures and the required data pipelines from unstructured data.

Faculty Frontier is strong in the decision stewardship use case, which involves applying governance principles to execute compliant, trusted and outcome-aligned decisions with accountability by monitoring quality standards and policies. Faculty Frontier's governance and decision monitoring capabilities help organizations ensure their AI-driven decisions adhere to regulatory requirements, ethical standards and organizational policies. This focus on compliance and accountability supports organizations in maintaining trust and transparency in their decision-making processes, which is critical for AI adoption. However, it currently lacks dedicated features for AI agents to assist decision stewards with monitoring and governance.

FICO

FICO Platform differentiates itself in decision engineering, offering developers and decision engineers an environment to compose, operationalize and scale decision

services using a modular, business-centric architecture. The platform features a low-code UI for decision analysts, supporting visual modeling of explainable decision models.

A major strength of FICO Platform is its advanced real-time event stream processing, which enables low-latency ingestion to process data streams and analysis of continuous data streams from diverse sources. The platform integrates multiple data streaming technologies, such as Kafka and its own native streaming, and is strong in complex event detection, with the ability to identify conditions, patterns and anomalies in event streams to trigger immediate responses. These include sending warnings, updating dashboards or initiating decision flows and workflows across business ecosystems. These capabilities are natively embedded for quick responses to real-time events. Another strength is FICO Platform's comprehensive library of decision blueprints and templates. These are curated by FICO and are reusable assets tailored to industries and individual business functions, including, but not limited to, banking and investment services. The templates, available via FICO Marketplace, define decision flows, data source connectivity, and personas and outcomes across the decision life cycle. The platform complements its optimization engine, FICO Platform - Enterprise Optimization, which applies prescriptive analytics and advanced algorithms to solve complex problems by maximizing or minimizing objectives under constraints.

The platform currently has limited support for AI agents for decision intelligence, natural language processing (NLP) and agentic DI capabilities for automating decision execution with LLMs. Further, the platform did not demonstrate SLM integration capabilities in general availability, although the FICO Focused Language Model was released in September 2025. Non-GenAI NLP techniques are part of FICO's separately marketed Customer Communication Service.

FICO Platform's strongest use case is decision engineering, where developers can compose and operationalize decision services for execution at scale, utilizing multimodal analytics preparation and data streams. Decision engineers will benefit from its decision service composition, particularly its composable architecture with business-centric software patterns. Its delivery approach provides decision engineers with a framework for modular software. It uses components that implement well-defined business capabilities, are encapsulated for autonomy, and are packaged for discovery and programmatic access.

FlexRule

FlexRule Open is a decision intelligence platform notable for its decision-centric approach, strong rule and logic modeling capabilities, and API-first architecture for decision service composition. The platform combines low-code authoring, declarative rule expression, and dynamic decision mapping to enable precise, adaptable decision workflows for both the decision analyst and the decision engineer, who will appreciate a choice between code-first, low-code or no-code in one unified platform.

FlexRule Open is strong in decision modeling, rule- and logic-based techniques, and decision service composition. Its decision-centric UI focuses on the decision workflow and allows in-context switching between UIs during design-time modeling. For decision logic, the platform offers a low-code user experience with various decision tables, FEEL boxed expressions and multiple languages for authoring decision models. Its strength in declarative rule expression is supported by accessible formats like decision tables and trees, enabling business users to define and articulate rule logic, as well as discover and generate rules. FlexRule Open's API-first approach makes decision services easy to access, use, govern and monitor. Developers benefit from this access method, which provides business software components that can be consumed and used by other applications, or composed with other components within the platform to embed decision services into business processes.

Areas for potential growth include graph and knowledge techniques, NLP, AI agents for decision intelligence, and real-time event stream techniques. FlexRule Open does not support graph analytics, which are important for understanding and exploring relationships between entities. It also lacks capabilities for AI agents in decision intelligence and execution, with no AI assistant available for use for monitoring or governance.

FlexRule Open's strongest use case is decision engineering, due to its capabilities in decision modeling. It is especially capable in defining mappings using the open standard Decision Model and Notation (DMN), Decision Requirement Diagrams and the vendor's Dynamic Decision Graph. Decision engineering is further enabled by strong decision service composition features for developers, thanks to its API-first design. FlexRule Open supports all evaluated use cases: decision analysis, decision engineering, decision science and decision stewardship. Decision-centric collaboration, however, remains challenging for nondevelopers and no-code users.

IBM

IBM offers Decision Intelligence, Cloud Pak for Data, and Cloud Pak for Business Automation, which include a range of products and services for decision intelligence, such as hybrid cloud, AI, software and IT infrastructure. It offers a versatile platform for large enterprises looking to advance their decision intelligence initiatives.

IBM is strong in machine learning, natural language techniques and decision service composition. IBM's ML capability allows it to extract insight and predict outcomes from high-dimensional datasets using statistical, computational and AI models. Techniques include supervised, semisupervised, unsupervised and reinforcement learning. IBM's platform(s) leverages advanced ML techniques to derive actionable insights, predict future outcomes and automate complex decision processes. This strength enables enterprises to address analytical needs, ranging from predictive modeling to optimization and anomaly detection. IBM also is strong in natural language techniques and decision service composition. IBM's platform provides users with the ability to componentize decision flows and encapsulate tasks using modular, reusable assets that are deployed as decision services, enabling organizations to build flexible, scalable and reusable decision logic. It also facilitates rapid development and deployment of decision-driven applications via its composable architecture. Its NLP capabilities enable users to extract insights from unstructured data and interact with decision systems using human language.

Areas for improvement for IBM are graph and knowledge techniques, real-time event stream techniques and decision collaboration. Its graph and knowledge techniques lack graph analytics and graph data science, which might inhibit exploration and discovery of relationships between entities and concepts such as organizations, people or transactions. Real-time event stream techniques in the platform(s) are weak in low-latency ingestion and continuous intelligence, reducing real-time awareness of contextual events.

IBM's strongest use case is decision science, which involves code-first use of the decision intelligence platform by decision scientists and operations researchers, who test experimental hypotheses by executing decisions to improve outcomes. This use case is enhanced by IBM's decision service composition, decision services and composable architecture, as well as strengths in ML to support advanced users who require sophisticated decision modeling, simulation and optimization tools to drive better business results.

InRule Technology

InRule Technology's InRule Decision Platform focuses on the development of reliable, trusted decision models that can be flexibly and easily deployed (via APIs) across multiple environments, including SaaS, public cloud (IaaS/PaaS), private cloud, hybrid cloud and on-premises configurations. InRule distinguishes itself by combining strong decision governance and explainable AI with modular decision service composition and flexible deployment options, enabling transparent, auditable decisioning at enterprise scale. Unlike platforms focused on advanced natural language, graph or agentic AI techniques, InRule prioritizes governance, compliance and operational agility for regulated industries through composite AI and reusable decision services.

The InRule Decision Platform is strong in decision governance, decision service composition and decision execution. Decision governance offers interactive visibility into the factors that drive ML models and allows explainable AI, including bias detection and mitigation capabilities. Decision service composition — enabled by encapsulation of decision rules, models and data connections as modular, reusable services — can be accessed and orchestrated via APIs. Those APIs ensure flexible deployment, governance and life cycle management of decision services. Decision execution supports real-time, batch and event-driven operations through DevOps integration (either within InRule or through Microsoft Azure or GitHub) using versioning, automated CI/CD pipelines, and other deployment options, including SaaS, cloud, hybrid, and on-premises configurations.

InRule has opportunities for improvement in natural language techniques, graph and knowledge techniques, and AI agents for decision intelligence. Natural language techniques are supported via REST API calls to third-party tools, such as OpenAI, but the platform lacks native integration of LLMs, SLMs and non-GenAI language processing. Graph and knowledge techniques are not offered, but users can extract information using the InRule SDK and logs to perform graph analytics in other tools. InRule does not support AI agents for DI, lacks agentic AI for decision model creation and collaboration for decision authoring, and does not support execution or monitoring agents.

InRule is best-suited for the decision stewardship use case because of its aforementioned strength in decision governance. It is also well suited for decision engineering because of its strengths in decision service composition and decision execution.

o9 Solutions

o9 Solutions, offers an AI-powered platform, the o9 Digital Brain platform, for integrated business planning and decision making for large enterprises. The o9 Digital Brain platform enables companies to integrate planning across supply chain, commercial and financial functions to optimize performance and align strategic objectives with operational execution.

o9's strongest capabilities are in ML, multimodal D&A preparation techniques and BI techniques. For ML, o9 can predict outcomes from datasets using statistical, computational and AI models that learn patterns and relationships. o9 Digital Brain leverages ML techniques to uncover actionable insights, predict future scenarios and augment decision flows, which is particularly valuable for supply chain optimization, demand forecasting and financial planning. o9 Solutions is strong in multimodal D&A preparation techniques, with data source connectivity allowing exploration and data ingestion across various sources. It also performs well in data preparation, providing a user-friendly interface to create data pipelines and analytical models. o9 also performs well in decision monitoring, allowing users to view each decision, model, logic and decision metadata across the decision flow. The platform's monitoring tools provide transparency and traceability that improve understanding of decisions and decision processes, and help to maintain accountability, ensure compliance and adapt to change.

o9 Solutions has an opportunity to improve decision collaboration, graph and knowledge techniques, and AI agents. For decision collaboration, its limited capabilities may require additional configuration to allow multiple users to work together across the DI life cycle. It could also improve human-AI delegation to balance human judgment and AI. The platform provides limited graph data science capabilities for decision scientists to predict relationships and patterns in connected data. o9 Solutions should consider enhancing the use of agentic AI in preproduction to collaboratively design and enable decision modeling and decision service composition.

o9 Solutions' strongest use case is decision engineering, aided by its support for decision services and composable architecture, as well as decision analytics for discovering, interpreting and communicating patterns in decision metadata to monitor decision quality, performance and outcomes. The platform provides code-first use by decision scientists to experiment with hypotheses to improve outcomes, demonstrating the ability to support advanced users requiring sophisticated modeling, simulation and optimization tools.

Oracle

Oracle's entry into the DIP market is Oracle Cloud Infrastructure (OCI), which includes Integration and Analytics Cloud. This offering builds on Oracle's comprehensive ecosystem capabilities to combine model-driven decisions and AI-enabled analytics that address deterministic decision needs. Oracle's DI platform consists of flexible OCI cloud services for D&A, ML and AI capabilities for data science, AI and GenAI, and integration and automation to embed decisions into workflows. The platform enables organizations to model, manage and automate complex decision logic using intuitive low-code tools. Integration of Oracle Analytics Cloud provides autogenerated insights, ML capabilities and AI-driven recommendations.

OCI is strong in multimodal D&A preparation techniques, ML techniques, decision monitoring and decision service composition. Comprehensive options across data source connectivity, data preparation and multistructured analytics solidify the strength of the multimodal D&A preparation capabilities, providing ease of use, smart recommendations and automated features (such as autojoins), which simplify the ability to easily and flexibly create data pipelines. For ML techniques, both supervised and unsupervised methods are readily accessible and easy to use for those needing the ability to leverage more advanced predictive analytics techniques. Decision observability is a key component of Oracle's decision monitoring capability, providing decision state indicators, detailed activity stream performance and metrics, visualization of dependencies, and business impact analysis. This enables transparency, auditability and explainability across the decision flow process. Decision service composition is built on Oracle's API-first approach and solid decision services capabilities.

Potential growth and areas for enhancement include AI agents for decision intelligence techniques, optimization/simulation techniques and decision collaboration. OCI currently does not support AI agents for DI techniques or optimization/simulation techniques within the platform. Decision collaboration is another area of opportunity, as the platform has only minimal support for decision-centric collaboration, human-AI delegation and embedded analytics within the decision flow.

OCI is best-suited for the decision engineering use case, with its strong multimodal D&A preparation techniques, real-time event stream techniques and decision observability. It is also well-suited for the decision science use case due to its strong decision observability as part of decision monitoring.

Pegasystems

Pegasystems provides Pega Infinity Platform, an offering that includes Customer Decision Hub, Process AI, Customer Experience Blueprint, Credit Risk Decisioning and Agentic Process Fabric. Customer Decision Hub provides customer engagement at scale and is core to the platform's decisioning capabilities. Other platform applications that focus on specific functions and processes include Pega Customer Service, Pega Sales Automation, Pega Customer Collections, Pega Process AI (for workflow decisions), Pega Credit Risk Management, Pega Care Management and Pega Claims Management for decision capabilities.

Pegasystems' key strengths include decision monitoring, multimodal D&A preparation techniques and real-time event stream techniques. Decision monitoring capabilities include strong decision analytics accessed via a decision observability dashboard. These capabilities assist with identifying bottlenecks in the decision flow and optimizing workflows for improved decision quality. Within multimodal D&A preparation techniques, Pegasystems provides consistent capabilities for multistructured analytics, data preparation and data source connectivity. Also included are real-time event stream techniques for low-latency ingestion and complex event detection, decision service composition supported via an API-first composable architecture, and decision collaboration capabilities.

Areas for improvement include graph and knowledge graphs, BI techniques and decision governance, all of which are lacking. Graph and knowledge graph capabilities can be improved across graph analytics, graph data science and knowledge graphs. These capabilities enable representation of complex relationships between entities and events across the decision flow. Similarly, BI techniques can be improved by strengthening the composite semantic layer to provide a consistent, governed view of the data across multiple sources, a multiexperience UI to accommodate varying levels of skill across roles and an action framework to help users compose and automate data-driven decision workflows. Finally, decision governance lacks strong capabilities for decision stewardship needed to enforce decision governance policies and standards, explainable AI to clarify a model's functioning to a specific audience, and AI trust, risk and security management to assure AI governance, trustworthiness, fairness, safety, reliability, security, privacy and data protection.

Pegasystems is best-suited for the decision engineering use case due, in large part, to its composable architecture, decision-centric UI and DecisionOps capabilities.

Quantexa

The Quantexa Decision Intelligence Platform delivers powerful entity resolution and graph analytics at scale that provide contextual information to support decisions based on relationships across structured and unstructured data sources. Quantexa focuses on decision quality, using the graph to provide a “Contextual Fabric” that acts as a composite semantic layer. The contextual fabric continuously integrates structured and unstructured data from internal and external sources, resolving entities and materializing relationships at scale to support decisions based on trends, hidden risks and relationships.

Quantexa Decision Intelligence Platform is strong in graph and knowledge techniques and decision governance. Graph and knowledge techniques sit at the core of Quantexa’s solution with patented technology to build and use “ego” (a graph generation technique) and knowledge graphs at scale. The platform identifies connections between entities and enables contextual insight for decision making. Decision governance is targeted at decision stewards by providing them with a dedicated interface for real-time monitoring and policy enforcement, including audit trail and model versioning, without engineering involvement.

Quantexa Decision Intelligence Platform can improve in optimization/simulation techniques, natural language techniques and AI agents for DI techniques.

Optimization/simulation techniques are not part of the Quantexa solution, although the ability to provide additional context, such as entity-level intelligence, makes it a useful preprocessing data feed to third-party optimization engines. Quantexa provides some capability to integrate natural language techniques, such as “bring-your-own” LLMs to query data, third-party fine-tuned SLMs and non-GenAI language processing techniques for a limited range of tasks such as data processing. Q Assist is used to support model design, execution and monitoring as AI agents for DI techniques are not currently available.

Quantexa is well-suited across the evaluated use cases, including decision engineering, decision stewardship and decision analysis. The decision engineering use case is strong because of the decision modeling and real-time event stream techniques. The decision stewardship use case stands out because of Quantexa’s strength in decision governance, modeling and real-time event stream techniques. The decision analysis use case also excels because of the platform’s strength in decision modeling. While graph and ML

capabilities enable the decision science use case, integration with other solutions may be needed to fulfill optimization, simulation, and rule- and logic-based techniques.

RelationalAI

The RelationalAI Decision Intelligence Platform provides a code-first development approach via notebooks and its “Rel” decision agent. The platform operates fully and natively within a data cloud platform, with Snowflake being the first (and currently only) supported option. The platform is open and interoperable, supporting open-table formats (such as Apache Iceberg) and open semantic interchange standards.

RelationalAI Decision Intelligence Platform’s strongest capabilities are graph and knowledge techniques, and machine learning techniques. Graph and knowledge techniques are particularly buoyed by the platform’s ability to provide graph data science and graph analytics via Python coding. This allows Python-proficient users to seamlessly apply graph techniques within decision intelligence workflows using familiar tools and environments. ML techniques are available via graph neural network (GNN) methods and are monitored and tracked through MLFlow. These include supervised learning, which builds models that learn from historical input-output pairs by training on labeled datasets, and unsupervised learning, which builds models that autonomously explore unlabeled data to uncover hidden patterns, clusters or associations without predefined output guidance. Semisupervised learning techniques that combine elements of both supervised and unsupervised learning are not available.

Decision collaboration, decision monitoring, decision governance and rule- and logic-based techniques are areas for potential growth and enhancement. For decision collaboration across the decision life cycle, human-AI delegation, embedded analytics and decision-centric collaboration are not sufficiently available. In addition, decision monitoring, decision analytics, learning from human feedback and decision observability are not adequately supported. Within decision governance, decision stewardship, explainable AI and AI trust, risk and security management are limited, with some capabilities accessible via Python coding. Within rule- and logic-based techniques, rule discovery and generation, declarative rule expression, and rule integrity and control are not supported. The components that comprise rule- and logic-based techniques are not available or limited by accessing a Python API.

Although offering limited support for the use cases defined in this analysis across the decision intelligence platform market, RelationalAI’s Decision Intelligence Platform

should be considered if an organization has a Snowflake ecosystem and prefers code-first environments.

Rulex

The Rulex Platform offers a no-code decision intelligence platform that provides data orchestration, explainable AI, scenario simulation and mathematical optimization to support and automate business decisions. The platform enables users to design decision workflows, harmonize data from multiple sources, and build tailored solutions through a visual interface, combining rule engines, AutoML and custom UIs. Rulex supports decision making and process automation across use cases such as fraud detection, inventory optimization and master data management in domains such as supply chain, finance and healthcare. Deployment options include on-premises, private or public cloud and containerized environments.

Rulex's most competitive capabilities are in machine learning techniques, multimodal D&A preparation and decision modeling. For ML, the platform offers a visual, no-code environment for training and deployment of both supervised and unsupervised models. It includes automated parameter selection for unsupervised approaches such as clustering and anomaly detection, allowing nonexperts to apply these methods. For multimodal D&A preparation, the platform delivers an easy-to-use, no-code interface for preparing and visualizing data from diverse sources and file types. Its strength in decision modeling comes from its customizable decision-centric UI, decision logic and composite AI.

Rulex has opportunities for improvement in decision governance, decision monitoring and real-time event stream capabilities. Its decision governance features offer limited support for decision stewardship, making it difficult for organizations to observe, manage and control decisions effectively. Although the platform provides visibility into the decision formula, it lacks significant capabilities to explain, secure and mitigate risks associated with AI-driven decisions, which can hinder trust and adoption. In decision monitoring, Rulex provides only basic and highly manual tools for decision observability. Rulex did not demonstrate analytics capabilities for monitoring or analyzing decision metadata, limiting the ability of users to continuously improve decision flows using the platform. Additionally, the platform does not support low-latency, real-time decisioning.

Compared with other use cases, Rulex is slightly better suited for decision analysis because of its decision modeling and decision logic capabilities, most notably its no-

code visual interface for modeling decision flows.

Sapiens

Sapiens' platform, Sapiens Decision, is built around "The Decision Model" (TDM) patented methodology. It provides essential elements of a user-friendly graphical model, with low-code authoring and graphical modeling of interconnected decisions. Its centralized repository provides reusable assets and robust governance capabilities.

Sapiens Decision is strong in rule- and logic-based techniques, decision modeling and decision governance. Rule- and logic-based techniques enable discovery and automated rules generation from various sources, including unstructured documents, structured rule environments and embedded rules in source code. Decision modeling, supported by the patented TDM methodology, enables users to graphically model, automate and govern decision logic using decision tables and diagrams. The user is assisted by AI-driven tools for logic discovery and optimization, and supported by a patented graphical model for showing and managing the relationship between decisions. Decision governance is enabled by placing the decision logic in a central repository using role-based access, customizable governance workflows and detailed audit trails. The platform incorporates explainability and interpretability and supports integration with ML/AI models, bias and fairness checks, and audience-specific reporting.

Sapiens Decision has the opportunity to improve in graph and knowledge techniques, natural language techniques and optimization/simulation techniques. For graph and knowledge techniques, it has no capabilities in graph analytics, graph data science or knowledge graphs. These areas, though underdeveloped, are critical for uncovering relationships, patterns and insights in connected data about decisions. Natural language techniques are limited to LLMs that support platform use. Although simulation techniques to build and test variations of decision models are available, support is limited for optimization and scenario planning techniques. AI agents for DI techniques are delivered via integration with third-party tools, with limited in-tool capability. However, the Sapiens Decision ModelAI tool is being extended with a Supervisor Agentic Pattern to develop tools that assist with information asset analysis, access and reuse.

Sapiens Decision is best-suited for the decision stewardship and decision engineering use cases. For decision stewards, the platform's strength is in decision governance, decision modeling and rule- and logic-based techniques. For decision engineers, the platform's strength is in decision modeling and decision service composition.

SAS

SAS Viya with Intelligent Decisioning is SAS Institute's product for decision intelligence. It is a comprehensive, end-to-end data, analytics, AI and decision intelligence platform, designed to support enterprise-scale analytics and decision making.

SAS Viya's strongest capabilities are multimodal data preparation, real-time event stream techniques and ML techniques that collectively provide the ability to handle complex, fast-changing and diverse data environments. SAS Viya's strength in multimodal data preparation techniques is critical and foundational to effective decision intelligence. It enables organizations to efficiently integrate, cleanse and transform diverse data sources, while supporting both analytics workflows and decision modeling. Its real-time event stream techniques provide the ability to process and analyze streaming data in real time offering low-latency ingestion, complex event detection and continuous intelligence to integrate time-sensitive data into operations. These techniques are suitable for industries and functions that require immediate insights and rapid response to contextual information, such as finance, healthcare and manufacturing. Additionally, the platform's advanced ML capabilities further enhance its value, by providing powerful tools for predictive and prescriptive modeling, pattern recognition and automated decision making.

SAS Viya has opportunities for improvement in graph and knowledge techniques, AI agents for DI techniques and BI techniques. While users can use a chat interface, the function is not integrated with the Viya Copilot. The model card in SAS Model Manager does not enable governance and monitoring via an agent out of the box. BI techniques need improvement, particularly when it comes to providing a composite semantic layer. The BI capability has yet to deploy semantic objects through the DI platform architecture for decision analysis.

The platform's strongest use case is decision engineering, providing developers with the ability to compose and operationalize decision services for execution at scale. Its strength in composite AI enables combined application of different AI techniques to improve the efficiency of learning and broaden the level of knowledge representations within the decision flow. SAS Viya enables developers to build, deploy and manage decision services that leverage its full, comprehensive spectrum of analytics, data preparation and AI capabilities.

Context

Decision intelligence platforms (DIPs) provide a solution that enhances how organizations make decisions, whether by humans or machines, individually or collectively. DIPs also address the growing challenge of making timely and accurate decisions for customers in disruptive, competitive and regulated markets. DIPs can create executable decision models that improve decision service composition and all-source intelligence to achieve better situational awareness, better recommendations or autonomous actions, tailored to specific decisions and outcomes. They can reduce the risk of poor decision making, allowing organizations to anticipate change and respond swiftly to opportunities at scale.

DIPs offer a coherent, contextual, connected, compliant, cost-optimized, collaborative and continuous approach to high-quality decision making in today's complex world. By optimizing the use of data, AI and other technology, they drive decision quality and better business outcomes (see [Innovation Insight: Decision Governance Mitigates Risks of Generative AI Agents](#).)

The precursor to the new Magic Quadrant for Decision Intelligence Platforms — [Market Guide for Decision Intelligence Platforms](#) — was first published on 18 July 2024. The 2026 [Magic Quadrant for Decision Intelligence Platforms](#) supersedes it. Our emerging Market Guide coverage was highly impactful and DIPs became a strategic area of focus within decision intelligence, data, analytics and AI practices. Client inquiries on DI topics rose steadily in 2024 and 2025.

Among many DI topics, Gartner is asked for assistance with end-user needs in today's market, such as the selection of technologies — including from DIP vendors — based on market position, maturity, best fit, strengths and weaknesses of their capabilities, known issues, market guidance and sizing, the technology landscape, and solution build-versus-buy platform and implementation (see [Toolkit: Decision Intelligence Platform Feature Checklist](#)).

A broader set of users, including no-code and low-code end users, are increasingly seeking solutions that provide a connected, contextual and continuous approach to decision making in today's complex environment. DIPs meet this demand by unifying and optimizing the use of data, AI and other technologies to enhance decision quality and drive better business outcomes.

Market Definition

Gartner defines decision intelligence platforms (DIPs) as software to create decision-centric solutions that support, augment and automate decision making of humans or machines, powered by the composition of data, analytics, knowledge and AI. DIPs enable enterprises to collaboratively design and explicitly model decisions, orchestrate decision flow during execution at scale, and enable monitoring and governance of decision quality, while learning from actions and outcomes. Features can include a combination of rule- and logic-based techniques, machine learning, real-time event stream processing, business intelligence, multimodal data and analytics preparation, natural language, graph technology, optimization, simulation or AI agents for decision intelligence.

DIPs provide a solution to enhance how organizations make decisions, whether by humans or machines, individually or collectively. They address the growing challenge of making timely and accurate decisions in volatile, uncertain, complex and ambiguous ecosystems, for more demanding customers in disruptive, competitive and regulated markets. DIPs help by creating executable decision models that improve decision service composition and all-source intelligence to achieve better situational awareness, better recommendations or autonomous actions, tailored to specific decisions and outcomes. They can reduce the risk of poor decisions, allow organizations to anticipate change and respond more swiftly to opportunities at scale.

DIPs help an organization achieve the following outcomes:

- **Improved decision making** — With the ability to model decisions and orchestrate decision flows, DIPs make the decision making life cycle structured and transparent, and decision services more adaptive based on learning from their environment, leading to better outcomes.
- **Increased efficiency** — Augmenting and automating decision making at scale saves time and resources, utilizing prescriptive analytics so teams can focus on more strategic tasks.
- **Enhanced collaboration** — Collaborative decision modeling and execution help ensure interdependent and diverse perspectives are considered, leading to more comprehensive and trusted decisions.

- **Better governance and accountability** — The ability to explain, evaluate, govern and audit decisions ensures organizations can learn from past actions and continuously improve trust, safety and compliance with regulations.

DIPs offer a coherent, contextual, connected, compliant, cost-optimized, collaborative, and continuous approach to high-quality decision making in today's complex world. By optimizing the use of data, AI and other technology, they drive decision quality and better business outcomes.

Mandatory Features

The mandatory features for this market are:

- **Decision modeling** — The capability to design explainable decision models from composite AI using a visual, low-code, decision-centric user interface to frame decisions with defined inputs, flow and outputs. It includes blueprints and decision network modeling to capture life cycle context, connectivity and continuity.
- **Decision collaboration** — The capability to improve human-AI delegation, minimizing friction between human and machine decision actors within teams, enterprises and ecosystems. It generates decision workflows, ethics and outcomes safeguards, guardrails and alerting thresholds for input/output monitoring and risk mitigation.
- **Decision service composition** — The capability to componentize decision flows and encapsulate tasks using modular, reusable components as decision services packaged for discovery. It includes integrating enterprise and partner system architectures and data ecosystems with an integration framework and programmatic access for composability.
- **Decision execution** — The capability to orchestrate and execute decision flows. It includes end-to-end decision life cycle management across development, testing and production environments, and various deployment options to ensure the reliable, scalable and efficient batch and real-time operations of decision services.
- **Decision monitoring** — The capability to view each decision, its model, logic, decision metadata and insight into the decision flow. It includes suggestions and alerts to implement safe and beneficial adaptations across the decision life cycle, involving humans in the loop as appropriate for learning and improving impact.

- **Decision governance** — The capability to apply governance principles to DI by logging, auditing and advancing decision making with an accountability framework for secure, safe, ethical, transparent, repeatable, outcome-led decisions; includes governing decisions as assets (model, logic, metadata): stewardship, policies and metrics.

Optional Features

The optional features for this market are:

- **Rule- and logic-based techniques** — The capability to capture organizational know-how, policies, procedures, standards, regulations and explicit knowledge as structured rules. Techniques include formal directives and expert heuristics or derivation from machine learning or optimization, controlled by analysts to ensure coherence.
- **Machine learning techniques** — The capability to extract insight and predict outcomes from high-dimensional datasets using statistical, computational and AI models that learn patterns and relationships. Techniques include supervised, semisupervised, unsupervised and reinforcement learning.
- **Real-time event stream techniques** — The capability to ingest data from streams with low latency for fast analytics and integration. Techniques include complex event detection and alerting to trigger or inform adaptive and continuous intelligence.
- **Business intelligence techniques** — The capability to model, analyze and visualize data via dashboards, reports and visualizations to support informed decision making. Techniques include empowering users to uncover patterns, predict trends, improve operations and collaborate on findings.
- **Multimodal D&A prep techniques** — The capability to unlock value from the spectrum of wide data in any format, language or image, sourced internally or externally. It deconstructs meaningful features of human-generated information for machines. It composes machine-generated data into context-enriched analysis for decisions.
- **Natural language techniques** — The capability to enable intuitive human-system communication by handling language aspects (pragmatic, semantic, grammatical, lexical). Techniques include generative AI and language processing methods to

extract rules or policies, and parse, interpret and generate language, translations and summaries.

- **Graph and knowledge techniques** — The capability to represent and analyze highly connected data, such as relationships between entities, network nodes or data objects. Techniques include graph data management, ontologies and semantic networks, knowledge graphs, graph analytics or causal graphs.
- **Optimization/simulation techniques** — The capability to maximize benefits and manage trade-offs by finding optimal resource combinations under constraints via operations research. Techniques include linear, constraint, integer programming, evolutionary algorithms, what-if analysis or scenario planning to simulate environments and outcomes.
- **AI agents for DI techniques** — The capability to use agentic AI to improve speed, quality and adaptability of decision making. It includes AI agents' (semi)autonomous role in decision modeling and composition, supporting, augmenting or automating decision service composition and execution, collaboration, monitoring and governance.

Product/Service Trends

End-user organizations express demand for DIPs due to their struggle with modeling decisions, improving efficiency through augmentation and automation, fostering collaboration, and ensuring strong governance. DIPs address these challenges by enabling improved decision modeling, streamlining processes, enhancing collaboration across teams, and strengthening governance and accountability — ultimately building trust, ensuring safety and supporting compliance with regulations.

End users are interested in the DIP market itself, and DI in the context of AI, generative AI and AI agents. They are also interested in monitoring decisions for evaluation of business outcomes and decision governance. Vendors have responded to these evolving end-user needs by differentiating product capabilities and marketing from their perceived competitors to stay competitive.

Our inclusion criteria for this research and companion Magic Quadrant reflect how the market is evolving as vendors balance two priorities: meeting cross-industry requirements while responding to strong demand for domain-specialized capabilities.

Vendors are expanding beyond foundational features to deliver preconfigured templates, low-code environments and composite AI techniques tailored for high-value use cases such as supply chain optimization and intelligent order fulfillment. At the same time, platforms are incorporating GenAI-driven decision flows, collaborative tools and modeling capabilities that support flexibility across multiple industries. This reflects a shift from purely data-driven approaches toward decision-centric architectures that combine deep domain expertise with configurable, industry-agnostic frameworks.

Critical Capabilities Definition

Decision Modeling

The capability to design explainable decision models from composite AI using a visual, low-code, decision-centric user interface to frame decisions with defined inputs, flow and outputs. It includes blueprints and decision network modeling to capture life cycle context, connectivity and continuity.

Decision Collaboration

The capability to improve human-AI delegation, minimizing friction between human and machine decision actors within teams, enterprises and ecosystems. It generates decision workflows, ethics and outcomes safeguards, guardrails and alerting thresholds for input/output monitoring and risk mitigation.

Decision Service Composition

The capability to componentize decision flows and encapsulate tasks using modular, reusable components as decision services packaged for discovery. It includes integrating enterprise and partner system architectures and data ecosystems with an integration framework and programmatic access for composability.

Decision Execution

The capability to orchestrate and execute decision flows. It includes end-to-end decision life cycle management across development, testing and production environments, and various deployment options to ensure the reliable, scalable and efficient batch and real-time operations of decision services.

Decision Monitoring

The capability to view each decision, its model, logic, decision metadata and insight into the decision flow. It includes suggestions and alerts to implement safe and beneficial adaptations across the decision life cycle, involving humans in the loop as appropriate for learning and improving impact.

Decision Governance

The capability to apply governance principles to DI by logging, auditing and advancing decision making with an accountability framework for secure, safe, ethical, transparent, repeatable, outcome-led decisions; includes governing decisions as assets (model, logic, metadata): stewardship, policies and metrics.

Rule- and Logic-Based Techniques

The capability to capture organizational know-how, policies, procedures, standards, regulations and explicit knowledge as structured rules. Techniques include formal directives and expert heuristics or derivation from machine learning or optimization, controlled by analysts to ensure coherence.

Machine Learning Techniques

The capability to extract insight and predict outcomes from high-dimensional datasets using statistical, computational and AI models that learn patterns and relationships. Techniques include supervised, semisupervised, unsupervised and reinforcement learning.

Real-Time Event Stream Techniques

The capability to ingest data from streams with low latency for fast analytics and integration. Techniques include complex event detection and alerting to trigger or inform adaptive and continuous intelligence.

Business Intelligence Techniques

The capability to model, analyze and visualize data via dashboards, reports and visualizations to support informed decision making. Techniques include empowering users to uncover patterns, predict trends, improve operations and collaborate on findings.

Multimodal D&A Prep Techniques

The capability to unlock value from the spectrum of wide data in any format, language or image, sourced internally or externally. It deconstructs meaningful features of human-generated information for machines and composes machine-generated data into context-enriched analysis for decisions.

Natural Language Techniques

The capability to enable intuitive human-system communication by handling language aspects (pragmatic, semantic, grammatical, lexical). Techniques include generative AI and language processing methods to extract rules or policies, and parse, interpret and generate language, translations and summaries.

Graph and Knowledge Techniques

The capability to represent and analyze highly connected data, such as relationships between entities, network nodes or data objects. Techniques include graph data management, ontologies and semantic networks, knowledge graphs, graph analytics or causal graphs.

Optimization/Simulation Techniques

The capability to maximize benefits and manage trade-offs by finding optimal resource combinations under constraints via operations research. Techniques include linear, constraint, integer programming, evolutionary algorithms, what-if analysis, scenario planning to simulate environments and outcomes.

AI Agents for DI Techniques

The capability to use agentic AI to improve speed, quality and adaptability of decision making. It includes AI agents' (semi)autonomous role in decision modeling and composition, supporting, augmenting or automating decision service composition and execution, collaboration, monitoring and governance.

Use Cases

Decision Analysis

No-code use of the DI platform by analysts responsible for decision modeling and monitoring while collaborating with decision makers and other DI practitioners.

Decision Engineering

Low-code use of the DI platform by developers to compose and operationalize decision services for execution at scale, with multimodal D&A prep, data streaming and AI agents.

Decision Science

Code-first use of the DI platform by specialists in decision science and operations research to test experimental hypotheses by executing decisions to improve outcomes.

Decision Stewardship

Use of the DI platform to apply governance principles to execute compliant, trusted, outcome-aligned decisions with accountability by monitoring quality standards and policies

Inclusion and Exclusion Criteria

This Critical Capabilities research uses the same inclusion criteria as its companion **Magic Quadrant for Decision Intelligence Platforms**. To qualify for inclusion in this research, vendors had to meet the following criteria:

1. **General Availability:** Product(s) evaluated should be generally available on 1 September 2025. General availability (GA) is defined as something a vendor's clients have in a production environment, rather than something they are testing or evaluating. GA doesn't make any assertions about quality, rather that the product is generally available for sale/use and is fully supported. Betas with limited distribution (e.g., invite only, limited user numbers) or finite duration (e.g., expiry date for use) are not generally available.
2. **Market Definition:** The product(s) must meet the Market Definition defined above.
3. **Cross-Industry:** The customer base for production deployment of the product since 1 January 2024 through 31 December 2024 must include paying customers for the product(s) in at least three of the following industries: Banking and Investment

Services; Communications, Media and Services; Education; Government; Healthcare; Insurance; Manufacturing and Natural Resources; Oil and Gas; Power and Utilities; Retail; Transportation; and Wholesale Trade. See **Note 1: Vertical Industries**.

4. **Customer Interest:** Vendor ranks among the Top 21 for the Customer Interest Indicator (CII) as defined by Gartner. CII was calculated using a weighted mix of internal and external inputs that reflect Gartner client interest, vendor customer engagement, and vendor customer sentiment between 1 January 2024 and 31 July 2025.

Weighting for Critical Capabilities in Use Cases

<i>Critical Capabilities</i>	<i>Decision Analysis</i>	<i>Decision Engineering</i>	<i>Decision Science</i>	<i>Decision Stewardship</i>
Decision Modeling	20%	15%	10%	10%
Decision Collaboration	10%	5%	5%	10%
Decision Service Composition	5%	20%	10%	5%
Decision Execution	5%	10%	5%	5%
Decision Monitoring	10%	10%	15%	15%
Decision Governance	5%	5%	5%	20%
Rule- and Logic-Based Techniques	10%	0%	10%	10%
Machine Learning	5%	5%	10%	0%

Source: Gartner (January 2026)

This methodology requires analysts to identify the critical capabilities for a class of products/services. Each capability is then weighed in terms of its relative importance for specific product/service use cases.

Critical Capabilities Rating - Data Provided by GICC

Each of the products/services that meet our inclusion criteria has been evaluated on a scale from 1.0 to 5.0 with 1 being the lowest rating and 5 being the highest.

Product/Service Rating on Critical Capabilities

<i>Critical Capabilities</i>	<i>ACTICO</i>	<i>Aera Technology</i>	<i>CRIF</i>	<i>Decisions</i>	<i>Faculty</i>
Decision Modeling	3.5	3.7	3.1	3.7	2.6
Decision Collaboration	3.2	2.5	2.2	3.6	2.4
Decision Service Composition	3.6	3.7	2.8	3.6	2.5
Decision Execution	3.7	2.8	2.0	3.7	2.1
Decision Monitoring	3.0	3.5	1.7	3.0	2.9
Decision Governance	3.7	2.9	1.8	2.8	2.2
Rule- and Logic-Based Techniques	2.7	2.9	2.4	3.5	1.7
Machine Learning	2.2	2.0	1.5	2.2	1.9

Source: Gartner (January 2026)

Table 3 shows the product/service scores for each use case. The scores, which are generated by multiplying the use-case weightings by the product/service ratings, summarize how well the critical capabilities are met for each use case.

Product Score in Use Cases

<i>Use Cases</i>	<i>ACTICO</i>	<i>Aera Technology</i>	<i>CRIF</i>	<i>Decisions</i>	<i>Faculty</i>	<i>FICO</i>
Decision Analysis	2.96	3.22	2.06	3.12	2.20	3.21
Decision Engineering	3.16	3.38	2.05	3.36	2.18	3.40
Decision Science	2.86	3.33	1.97	2.95	2.22	3.26
Decision Stewardship	3.15	3.10	1.98	3.28	2.24	3.22
As of 8 December 2025						

Source: Gartner (January 2026)

To determine an overall score for each product/service in the use cases, multiply the ratings in Table 2 by the weightings shown in Table 1.

Acronym Key and Glossary Terms

AI	Artificial intelligence
BI	Business intelligence
D&A	Data and analytics
DI	Decision intelligence
DIP	Decision intelligence platform

FEEL	Friendly Enough Expression Language
GenAI	Generative artificial intelligence
LLM	Large language model
ML	Machine learning
NLP	Natural language processing
SLM	Small language model

Note 1: Vertical Industries

Our Critical Capabilities research focuses only on the 12 primary or Tier 1 vertical industries, as defined in [Market Definitions and Methodology: Vertical Industries](#). This guide explains Gartner's vertical industry market segmentation, the definitions therein, and the research methodologies and metrics that form the basis of our vertical industry market research and enterprise IT spend forecast statistics.

Critical Capabilities Methodology

This methodology requires analysts to identify the critical capabilities for a class of products or services. Each capability is then weighted in terms of its relative importance for specific product or service use cases. Next, products/services are rated in terms of how well they achieve each of the critical capabilities. A score that summarizes how well they meet the critical capabilities for each use case is then calculated for each product/service.

"Critical capabilities" are attributes that differentiate products/services in a class in terms of their quality and performance. Gartner recommends that users consider the set of critical capabilities as some of the most important criteria for acquisition decisions.

In defining the product/service category for evaluation, the analyst first identifies the leading uses for the products/services in this market. What needs are end-users looking to fulfill, when considering products/services in this market? Use cases should match common client deployment scenarios. These distinct client scenarios define the Use Cases.

The analyst then identifies the critical capabilities. These capabilities are generalized groups of features commonly required by this class of products/services. Each capability is assigned a level of importance in fulfilling that particular need; some sets of features are more important than others, depending on the use case being evaluated.

Each vendor's product or service is evaluated in terms of how well it delivers each capability, on a five-point scale. These ratings are displayed side-by-side for all vendors, allowing easy comparisons between the different sets of features.

Ratings and summary scores range from 1.0 to 5.0:

1 = Poor or Absent: most or all defined requirements for a capability are not achieved

2 = Fair: some requirements are not achieved

3 = Good: meets requirements

4 = Excellent: meets or exceeds some requirements

5 = Outstanding: significantly exceeds requirements

To determine an overall score for each product in the use cases, the product ratings are multiplied by the weightings to come up with the product score in use cases.

The critical capabilities Gartner has selected do not represent all capabilities for any product; therefore, may not represent those most important for a specific use situation or business objective. Clients should use a critical capabilities analysis as one of several sources of input about a product before making a product/service decision.

Unlock more AI value

With unmatched, proprietary AI insights from Gartner

[Explore the AI Hub ↗](#)

© 2026 Gartner, Inc. and/or its affiliates. All rights reserved. Gartner is a registered trademark of Gartner, Inc. and its affiliates. This publication may not be reproduced or distributed in any form without Gartner's prior written permission. It consists of the opinions of Gartner's research organization, which should not be construed as statements of fact. While the information contained in this publication has been obtained from sources believed to be reliable, Gartner disclaims all warranties as to the accuracy, completeness or adequacy of such information. Although Gartner research may address legal and financial issues, Gartner does not provide legal or investment advice and its research should not be construed or used as such. Your access and use of this publication are governed by [Gartner's Usage Policy](#). Gartner prides itself on its reputation for independence and objectivity. Its research is produced independently by its research organization without input or influence from any third party. For further information, see "[Guiding Principles on Independence and Objectivity](#)." Gartner research may not be used as input into or for the training or development of generative artificial intelligence, machine learning, algorithms, software, or related technologies.

[About](#) [Careers](#) [Newsroom](#) [Policies](#) [Site Index](#) [IT Glossary](#) [Gartner Blog](#)
[Network](#) [Contact](#) [Send Feedback](#)

© 2026 Gartner, Inc. and/or its Affiliates. All Rights Reserved.